

# 211: Computer Architecture Spring 2017

Instructor: Prof. Santosh Nagarakatte

## Topics:

- C programming conclusion
- Data Representation
  - Reading: Chapter 2.1, 2.2, 2.3, 2.4, 2.5

# Signed/Unsigned in C

- Unsigned values in C
  - Declare unsigned int i = 10;
  - Use typecasts i = (unsigned) j;
- -1 when interpreted as a unsigned value is a huge number is
  - $2^n - 1$

# Sign extension

Signed integer	4-bit representation	8-bit representation	16-bit representation
+1	0001	00000001	0000000000000001
-1	1111	11111111	1111111111111111

# Floating point

Integers typically written in ordinary decimal form

- E.g., 1, 10, 100, 1000, 10000, 12456897, etc.

But, can also be written in scientific notation

- E.g.,  $1 \times 10^4$ ,  $1.2456897 \times 10^7$

What about binary numbers?

- Works the same way
- $0b100 = 0b1 \times 2^2$

Scientific notation gives a natural way for thinking about floating point numbers

- $0.25 = 2.5 \times 10^{-1} = 0b1 \times 2^{-2}$

How to represent in computers? Why not use fixed point numbers?

# IEEE floating point standard

Most computers follow IEEE 754 standard

Single precision (32 bits)

Double precision (64 bits)

Extended precision (80 bits)



# Numerical Values

## Three different cases:

### ■ Normalized values

- exponent field  $\neq 0$  and exponent field  $\neq 2^k - 1$  (all 1's)
- exponent = binary value – Bias
  - » Bias =  $2^{k-1} - 1$  (e.g., 127 for float)
- Value of the number =  $1.(\text{mantissa field})$
- Ex: (sign: 0, exp: 1, mantissa: 1) would give  $0b1.1 \times 2^{-126}$

### ■ Denormalized values

- exponent field = 0
- exponent =  $1 - \text{Bias}$  (e.g., -126 for float)
- Value of the number = mantissa field (no leading 1)
- Ex: (sign: 0, exp: 0, mantissa: 10) would give  $0b10 \times 2^{-126}$

### ■ Special values: represent $+\infty$ , $-\infty$ , and NaN

# Decimal to IEEE Floating Point

5.625

In binary

$101.101 \rightarrow 1.01101 \times 2^2$

Exponent field has value 2

- add 127 to get 129

Exponent is 10000001

Mantissa is 01101

Sign bit is 0

0 10000001 011010000000000000000000

# Floating point in C

## 32 bits single precision (type float)

- 1 bit for sign, 8 bits for exponent, 23 bits for mantissa
  - Sign bit: 1 = negative numbers, 0 = positive numbers
  - Exponent is power of 2
- Have 2 zero's
- Range is approximately  $-10^{38}$  to  $10^{38}$

## 64 bits double precision (type double)

- 1 bit for sign, 11 bits for exponent, 52 bits for mantissa
- Majority of new bits for mantissa → higher precision
- Range is  $-10^{308}$  to  $+10^{308}$



# One more example

Convert 12.375 to floating point representation

Binary is 1100.011

$$1.100011 \times 2^3$$

$$\text{Exponent} = 127 + 3 = 130 = 0b10000010$$

$$\text{Mantissa} = 100011$$

$$\text{Sign} = 0$$

# Floating Point Operations

- No exact representation for a floating point
  - Mantissa is only 23 bits in 32 bit representation
  - Least significant bits may be dropped
- Floating point operations are not associative
  - $(3.14 + 1e10) - 1e10 \neq 3.14 + (1e10 - 1e10)$ . Why?

# iClicker Pop Quiz 1

Let's say we have a 6-bit floating point representation with 1-bit for the sign, 3-bits for the exponent and 2-bits for the Mantissa.

What is the bias?

A: 7

B: 3

C: 1

D: 8

## iClicker Pop Quiz 2

Let's say we have a 6-bit floating point representation with 1-bit for the sign, 3-bits for the exponent and 2-bits for the Mantissa.

What is the abstract representation for a normalized value?

A:  $(-1)^S \cdot (1.M) \cdot 2^{E-7}$

B:  $(-1)^S \cdot (1.M) \cdot 2^{E-3}$

C:  $(1.M) \cdot 2^{E-3}$

D:  $2^{E-3}$

# iClicker Pop Quiz 3

Let's say we have a 6-bit floating point representation with 1-bit for the sign, 3-bits for the exponent and 2-bits for the Mantissa.

Which of this is a normalized floating point value?

A: 1 001 11

B: 0 111 00

C: 1 110 11

D: 0 000 11

# iClicker Pop Quiz 4

Let's say we have a 6-bit floating point representation with 1-bit for the sign, 3-bits for the exponent and 2-bits for the Mantissa.

Which of this is a denormalized floating point value?

A: 0 001 11

B: 1 000 11

C: 0 100 00

D: 1 111 11

# iClicker Pop Quiz 5

Let's say we have a 6-bit floating point representation with 1-bit for the sign, 3-bits for the exponent and 2-bits for the Mantissa.

Which of this is a representation of infinity?

A: 1 101 11

B: 0 000 11

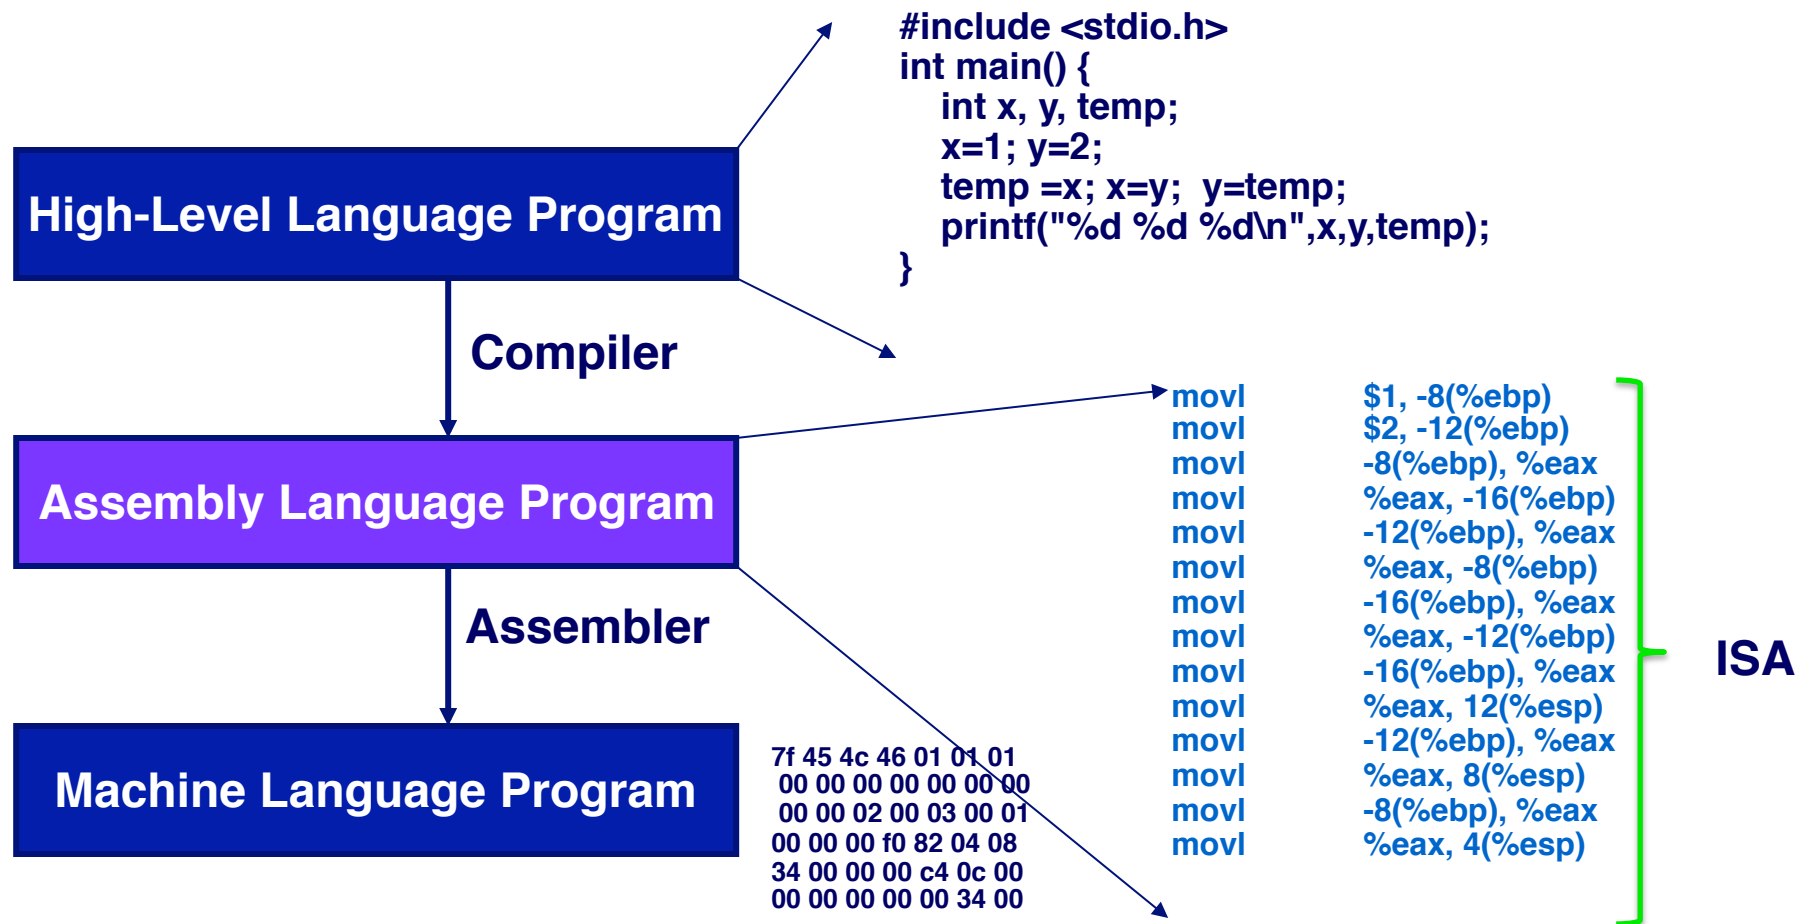
C: 0 111 11

D: 0 111 00

# Hardware Software Interface



# Programming Meets Hardware



How do you get performance?

# Performance with Programs

(1) Program: Data structures + algorithms

(2) Compiler translates code

(3) Instruction set architecture

(4) Hardware Implementation

# Instruction Set Architecture

- (1) Set of instructions that the CPU can execute
  - (1) What instructions are available?
  - (2) How the instructions are encoded? Eventually everything is binary.
- (2) State of the system (Registers + memory state + program counter)
  - (1) What instruction is going to execute next
  - (2) How many registers? Width of each register?
  - (3) How do we specify memory addresses?
    - Addressing modes
- (3) Effect of instruction on the state of the system

# IA32 (X86 ISA)

There are many different assembly languages because they are processor-specific

- IA32 (x86)
  - x86-64 for new 64-bit processors
  - IA-64 radically different for Itanium processors
  - Backward compatibility: instructions added with time
- PowerPC
- MIPS

We will focus on IA32/x86-64 because you can generate and run on iLab machines (as well as your own PC/laptop)

- IA32 is also dominant in the market although smart phone, eBook readers, etc. are changing this

# X86 Evolution

8086 – 1978 – 29K transistors – 5-10MHz

I386 – 1985 – 275K transistors – 16-33 MHz

Pentium4 – 2005 – 230M transistors – 2800-3800 MHz

Haswell – 2013 – > 2B transistors – 3200-3900 MHz

## Added features

- Large caches
- Multiple cores
- Support for data parallelism (SIMD) eg AVX extensions

# CISC vs RISC

CISC: complex instructions : eg X86

- Instructions such as strcpy/AES and others
- Reduces code size
- Hardware implementation complex?

RISC: simple instructions: eg Alpha

- Instructions are simple add/ld/st
- Increases code size
- Hardware implementation simple?

# Aside About Implementation of x86

About 30 years ago, the instruction set actually reflected the processor hardware

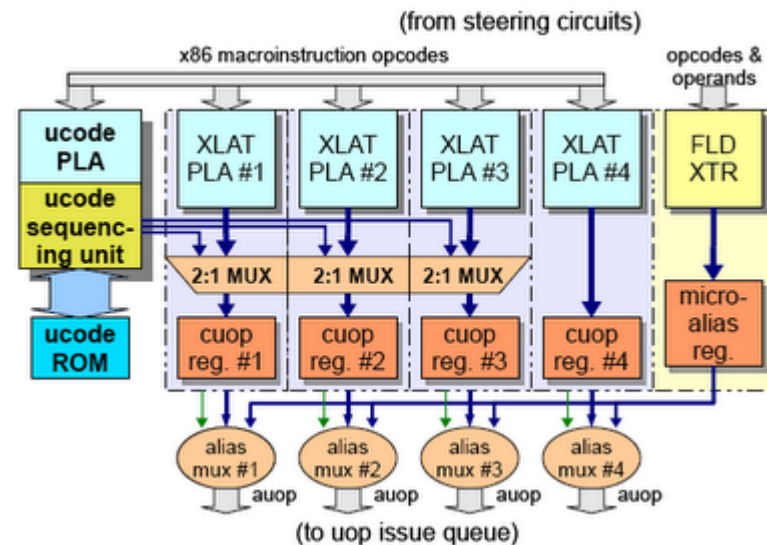
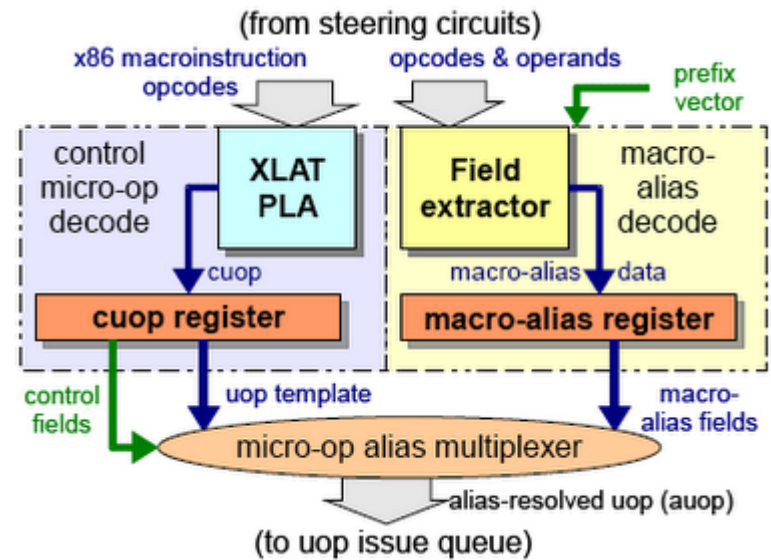
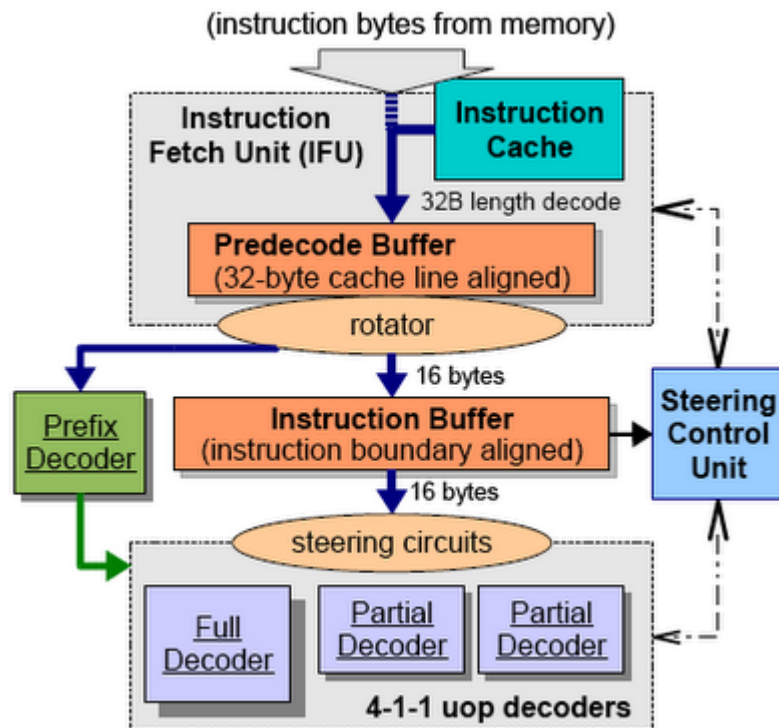
- E.g., the set of registers in the instruction set is actually what was present in the processor

As hardware advanced, industry faced with choice

- Change the instruction set: bad for backward compatibility
- Keep the instruction set: harder to exploit hardware advances
  - Example: many more registers but only small set introduced circa 1980

Starting with the P6 (PentiumPro), IA32 actually got implemented by Intel using an “interpreter” that translates IA32 instructions into a simpler “micro” instruction set

# P6 Decoder/Interpreter





# Assembly Programming

Brief tour through assembly language programming

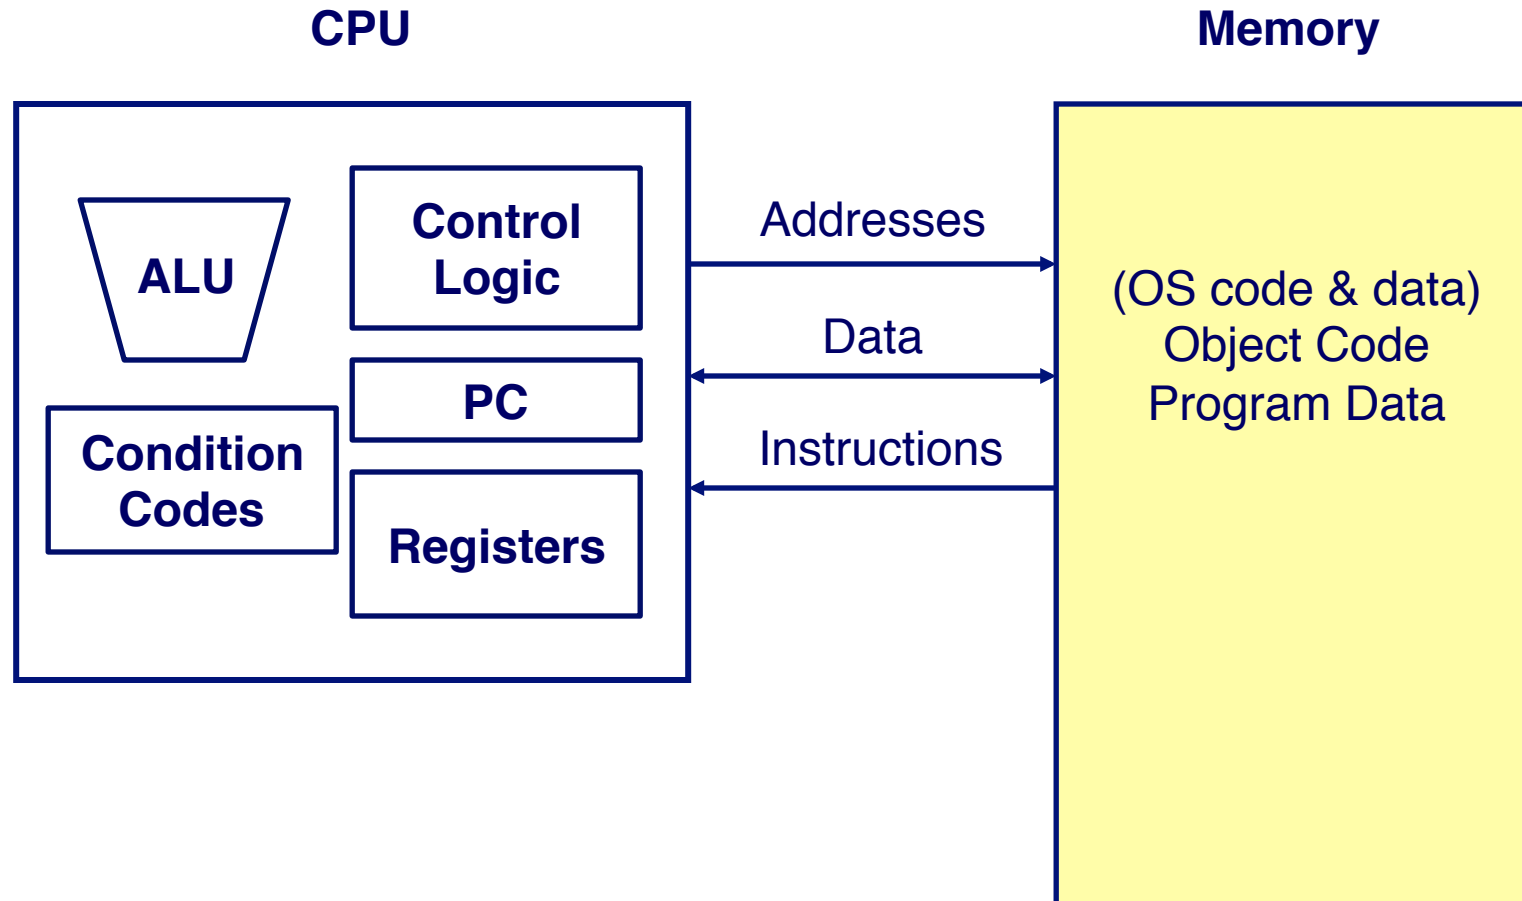
Why?

- Machine interface: where software meets hardware
- To understand how the hardware works, we have to understand the interface that it exports

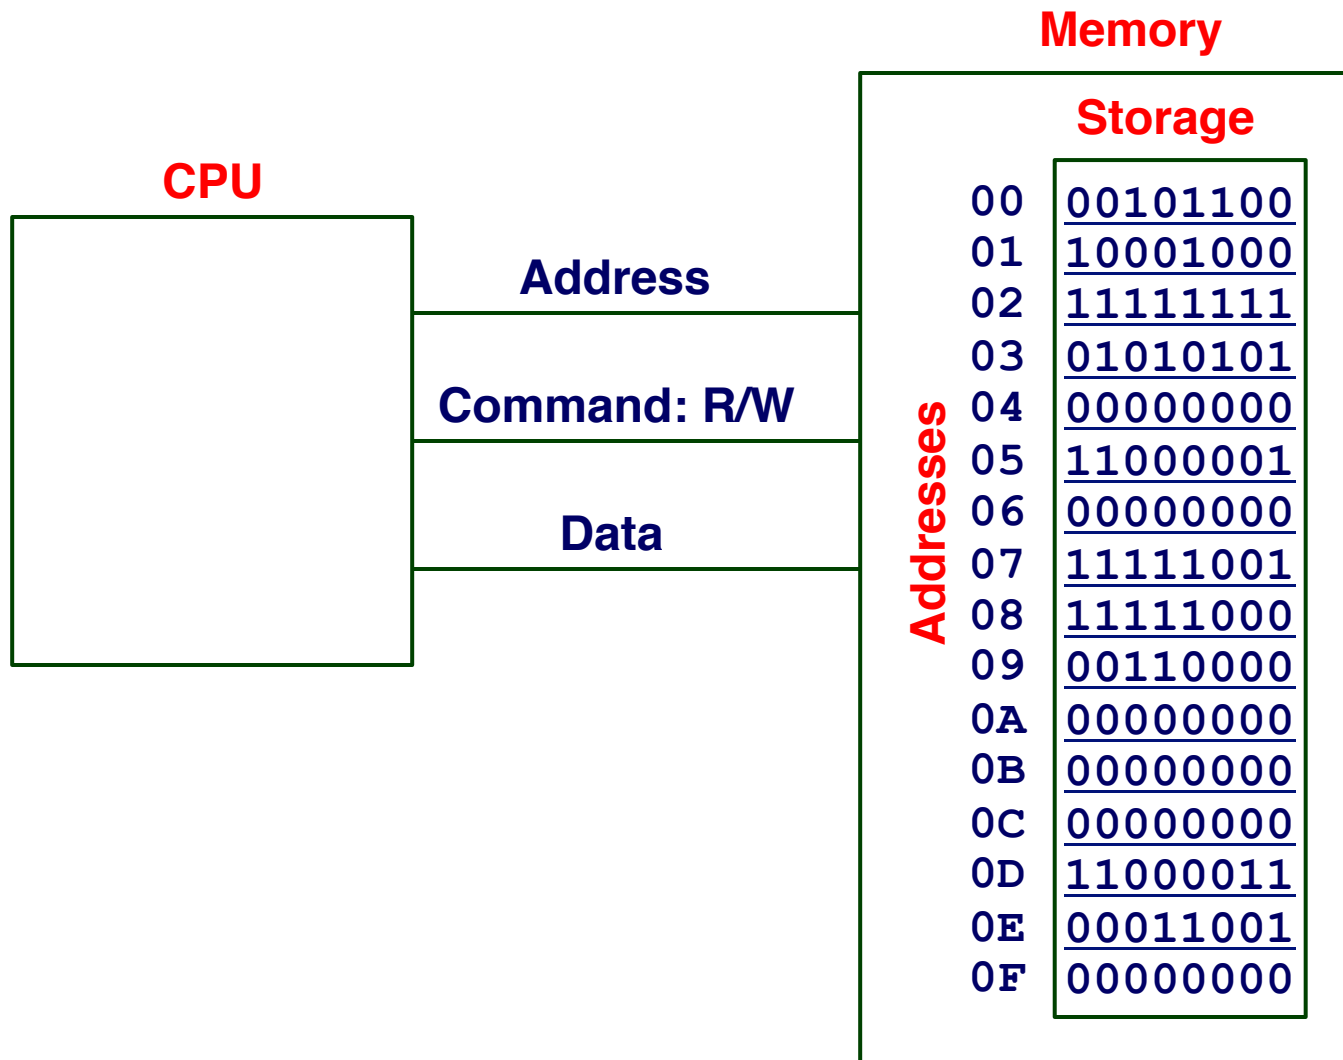
Why not binary language?

- Much easier for humans to read and reason about
- Major differences:
  - Human readable language instead of binary sequences
  - Relative instead of absolute addresses

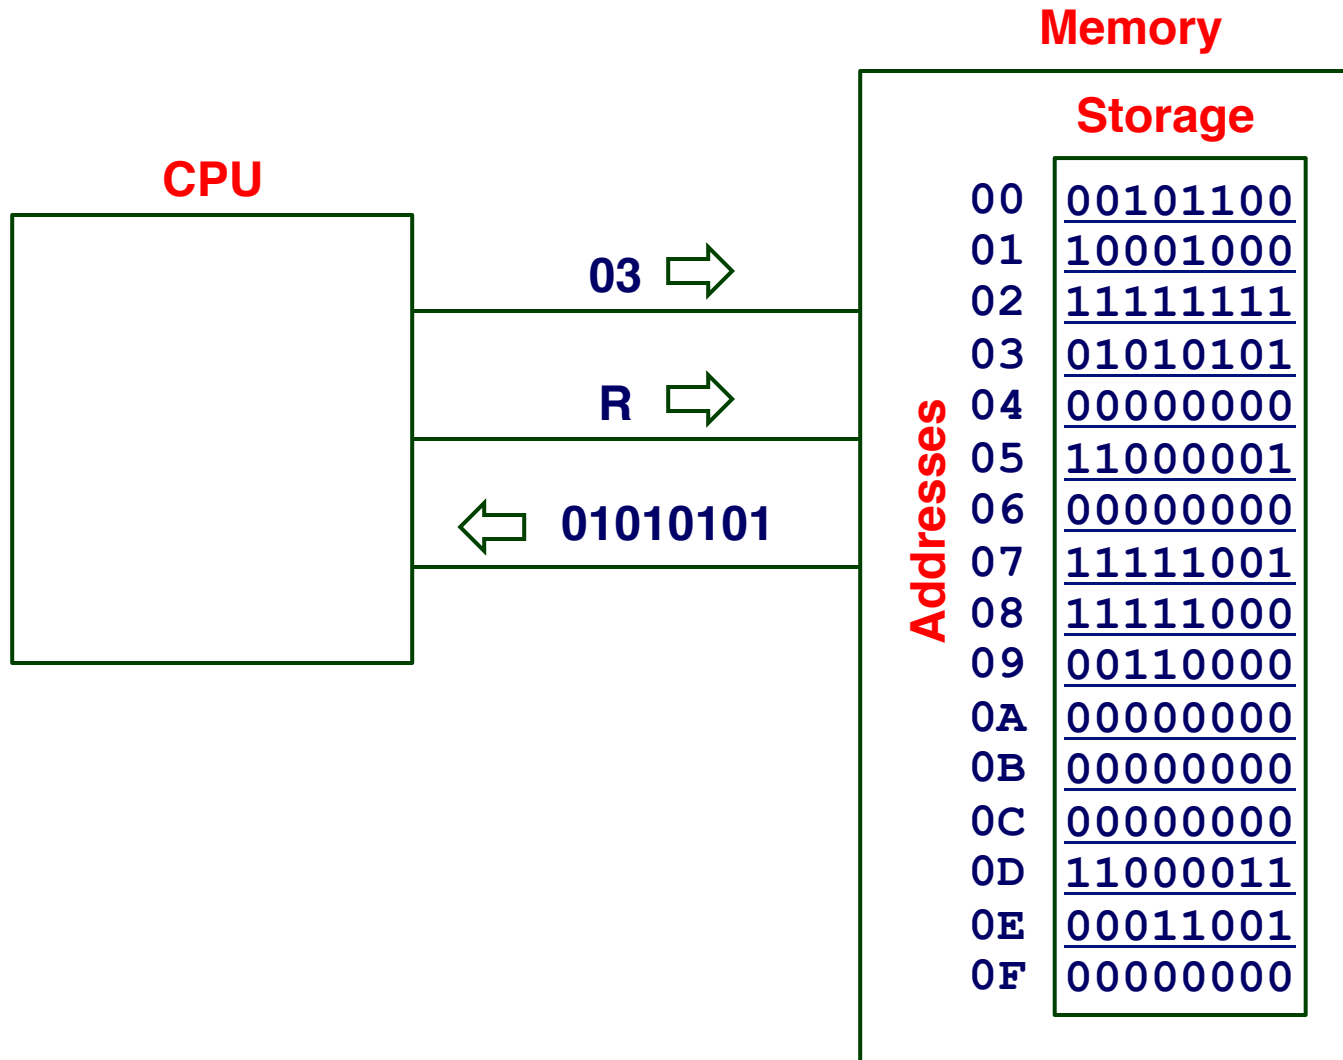
# Assembly Programmer's View



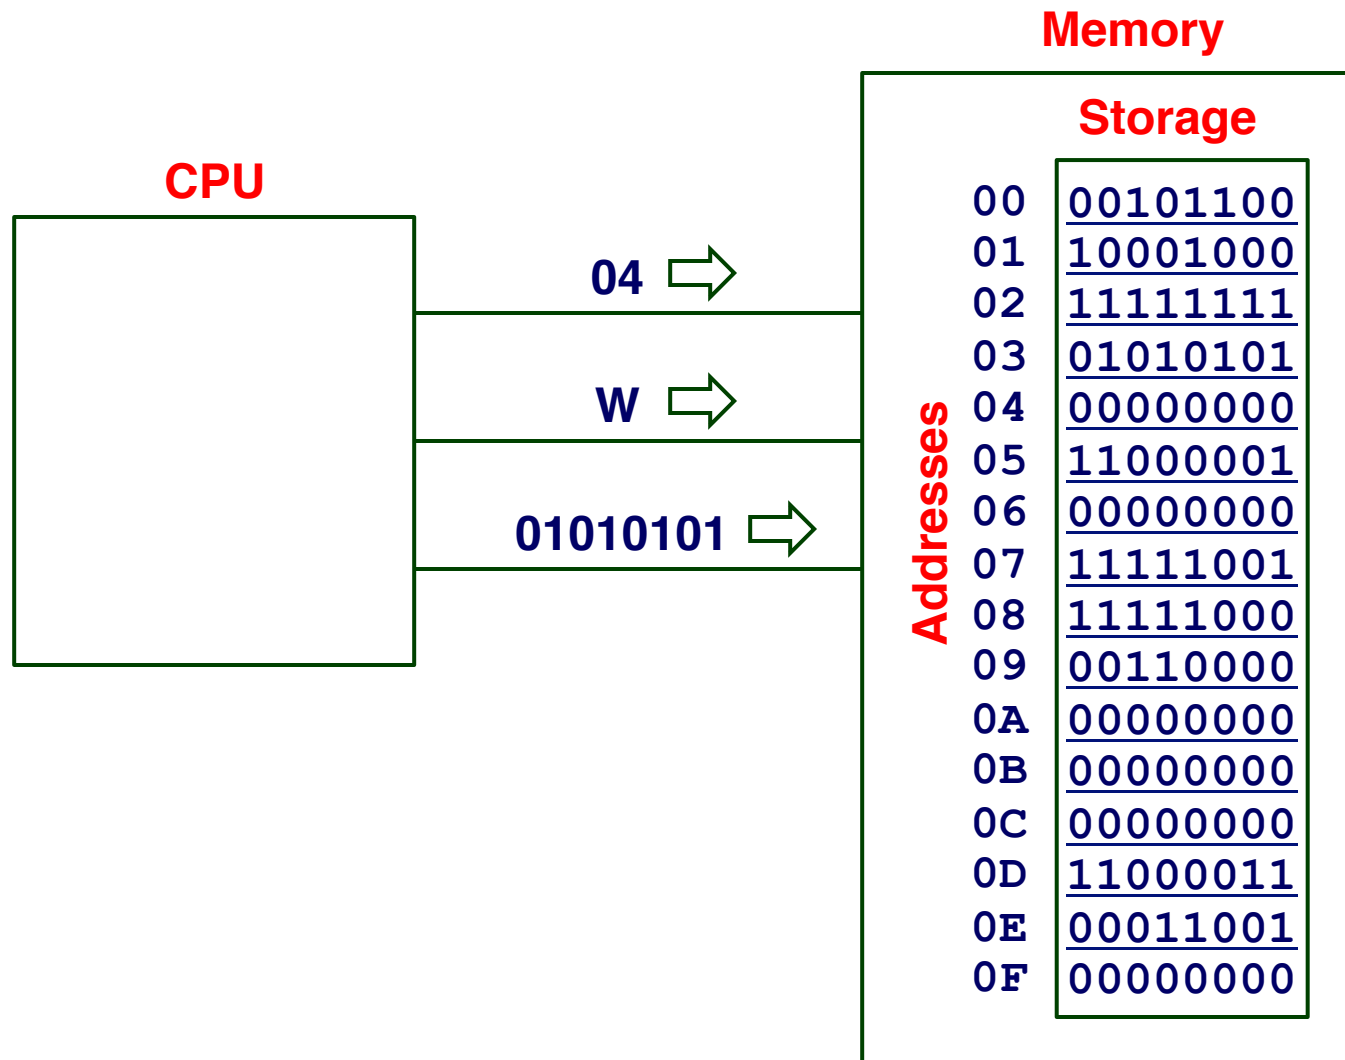
# Memory



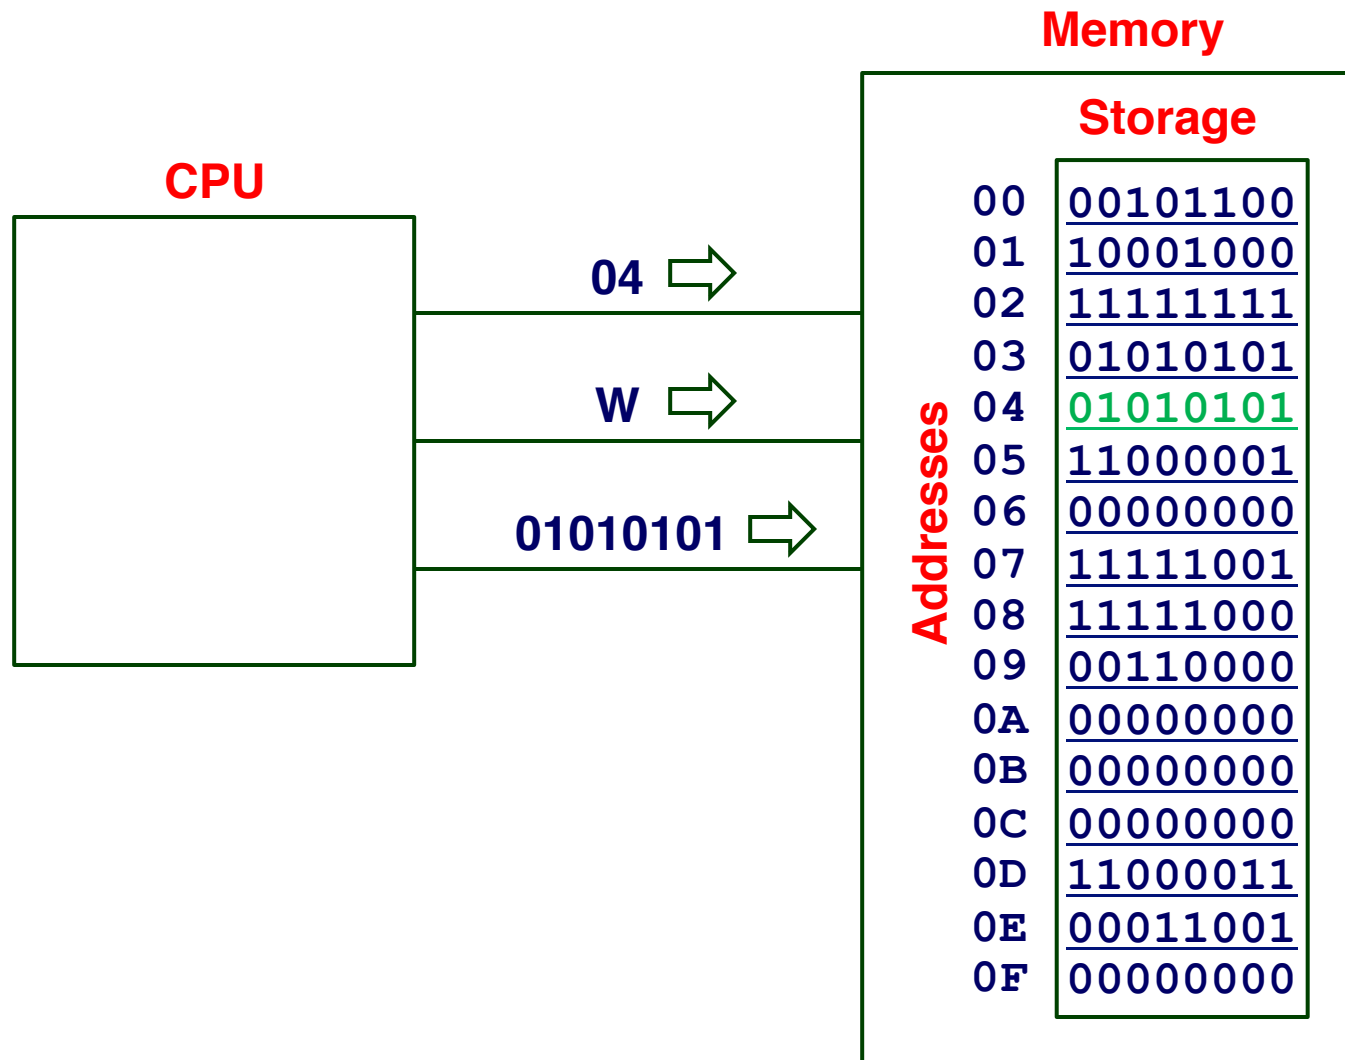
# Memory Access: Read




# Memory Access: Write



# Memory Access: Write

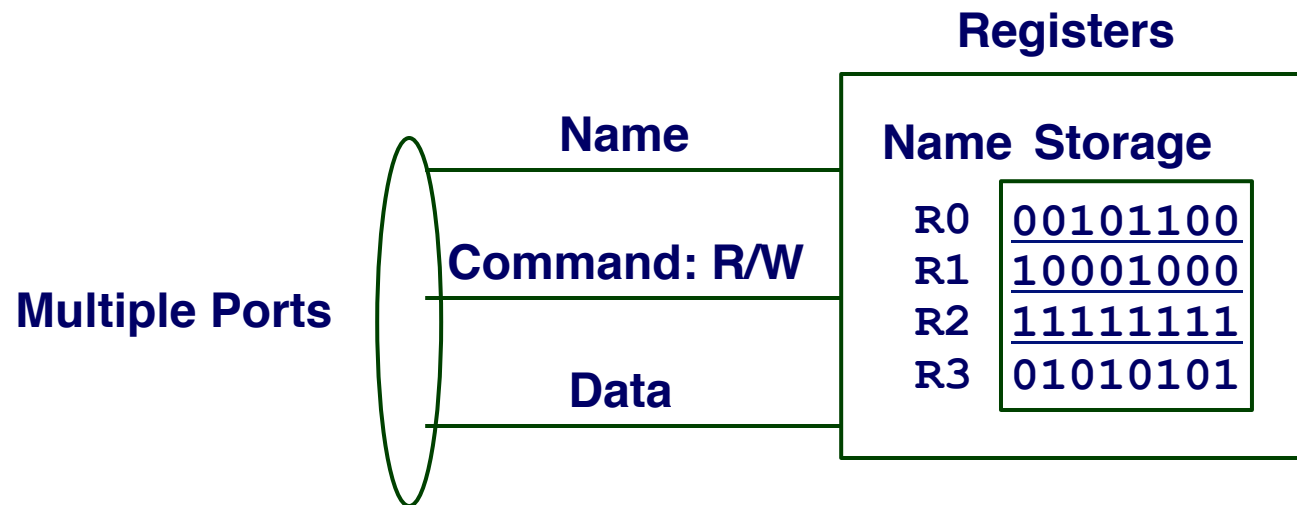


# Processor: ALU & Registers

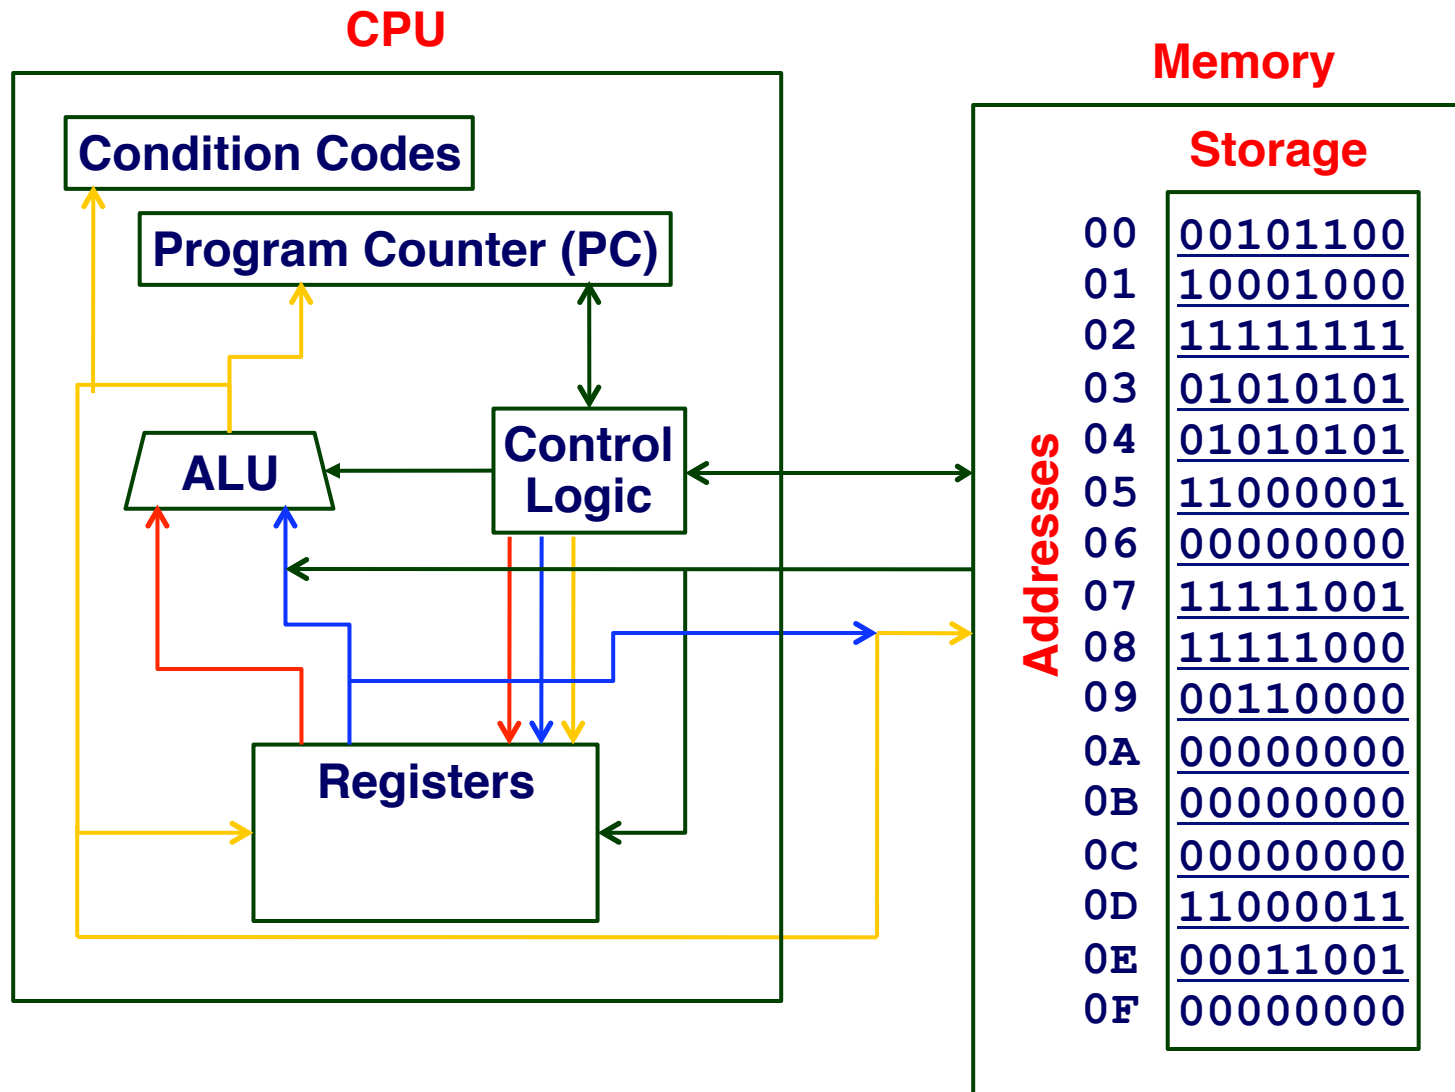


$$C = F_S(A, B)$$

F includes  
Arithmetic: +, -, \*, /, ~, etc.  
Logical: <, >, =, etc.

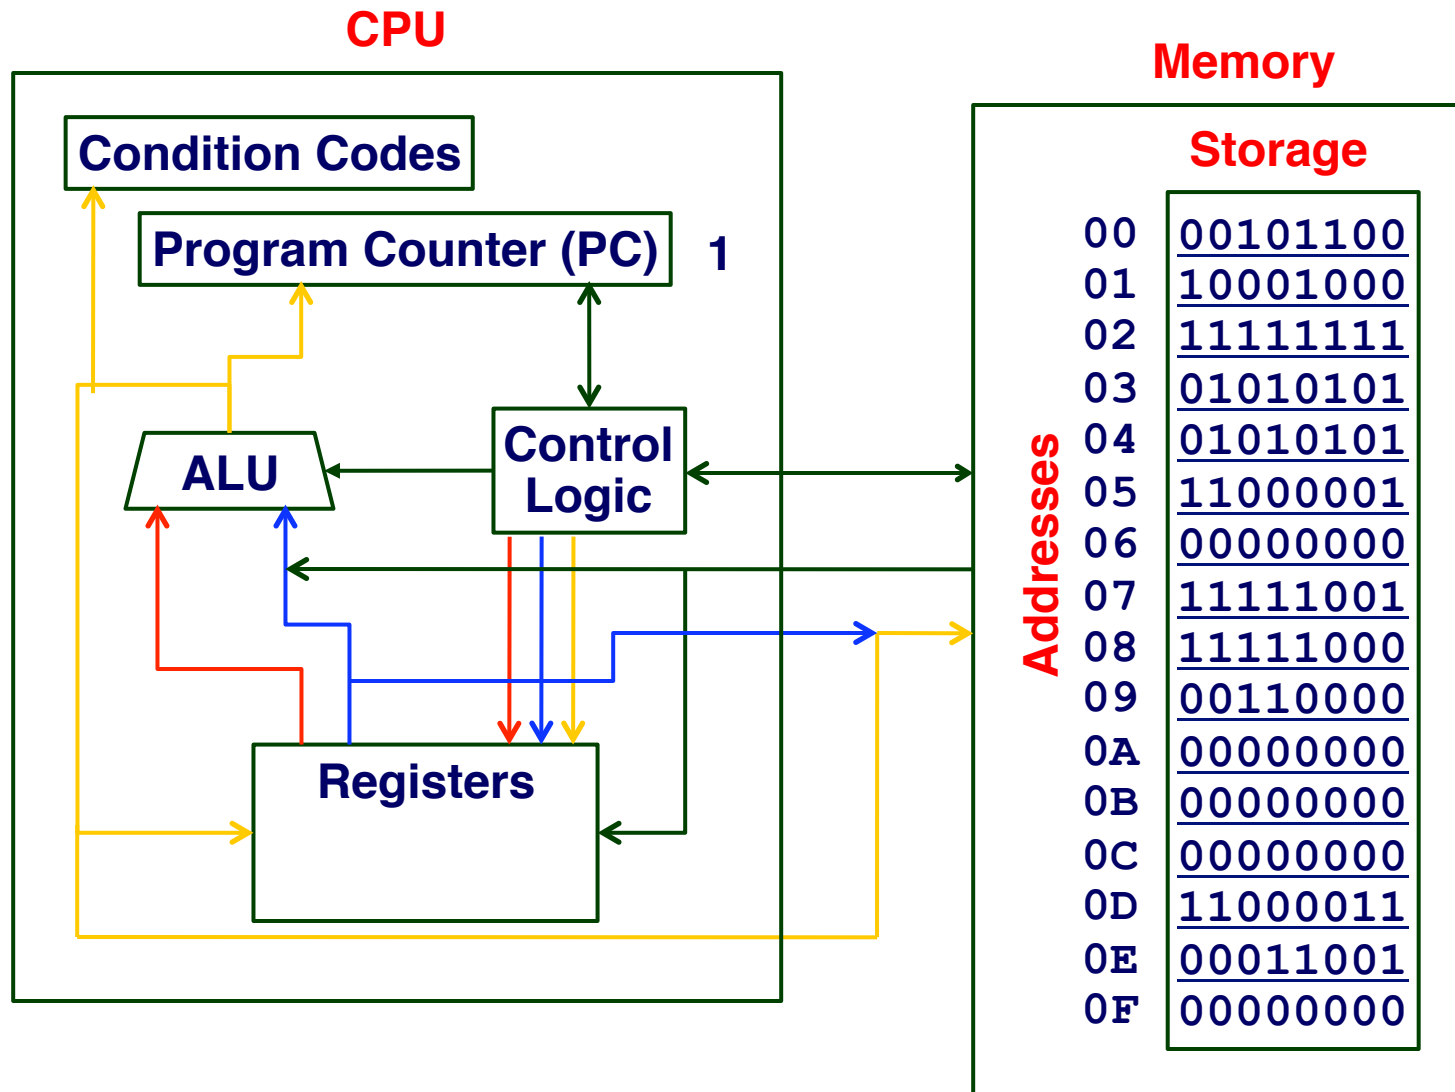


# Putting It All Together

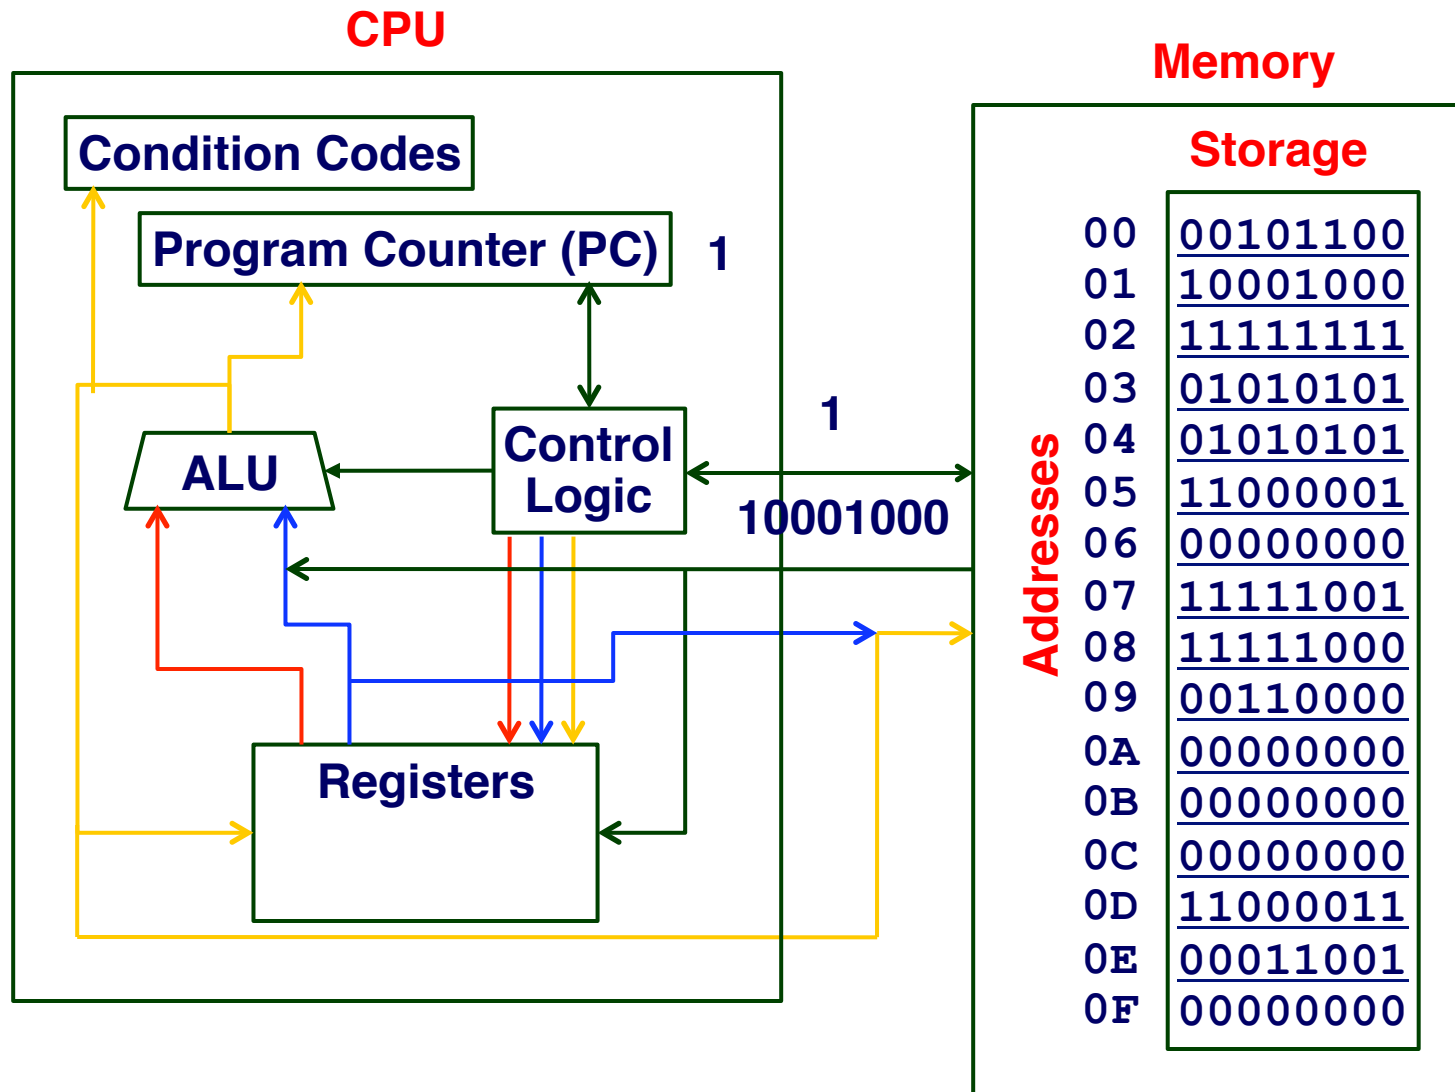




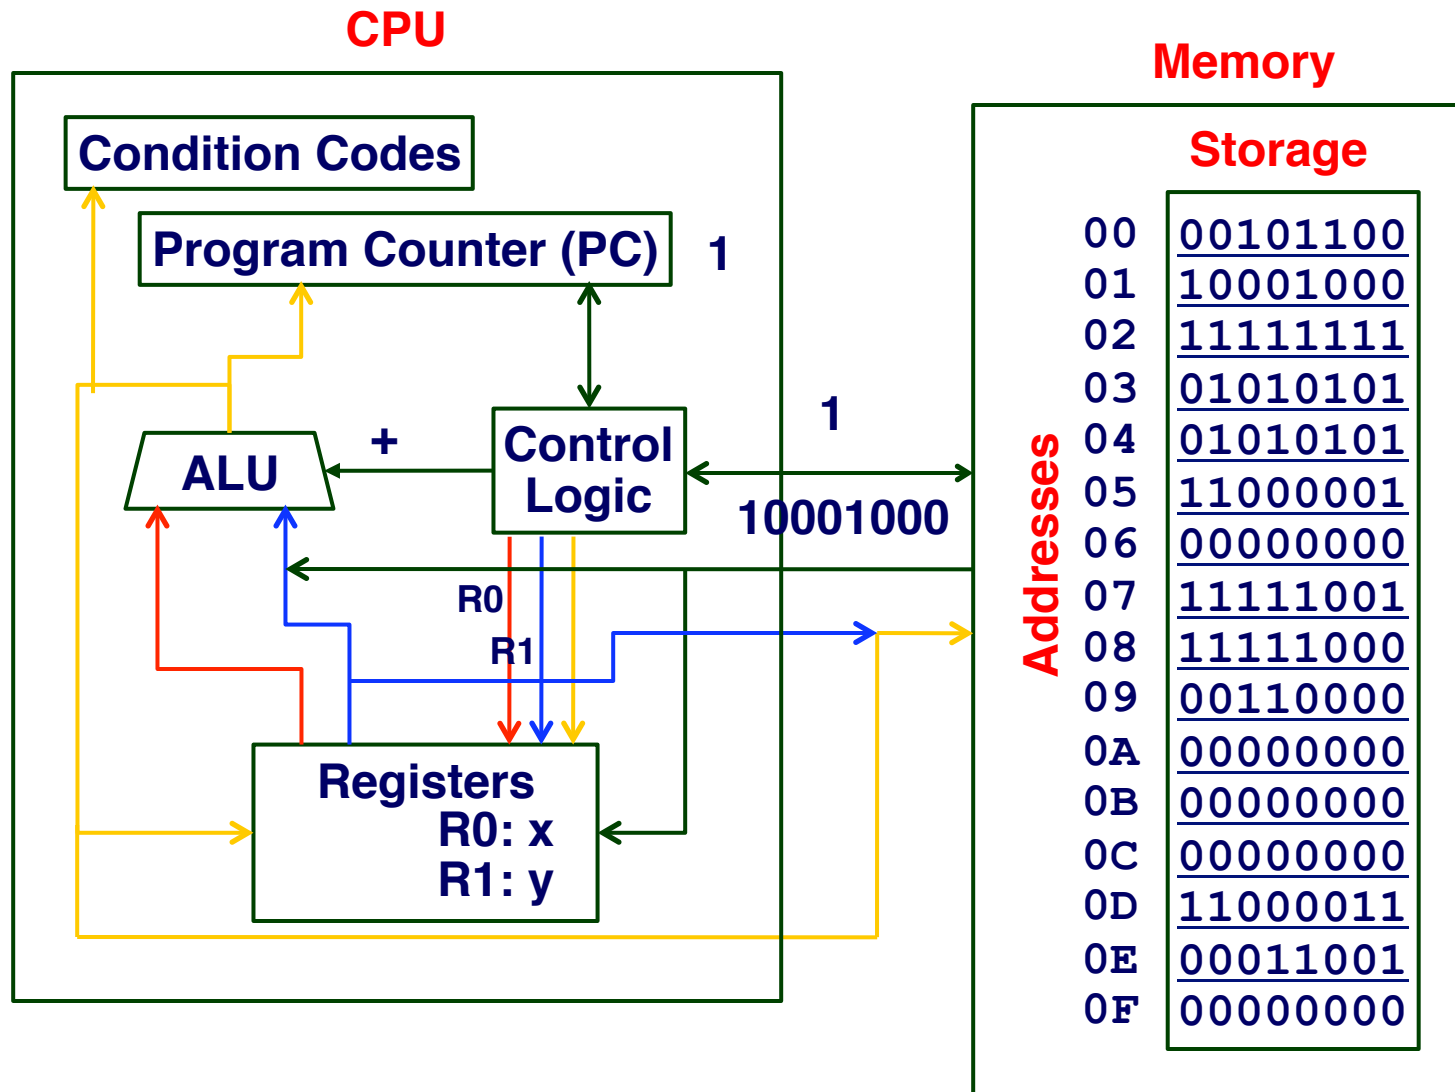
# Putting It All Together



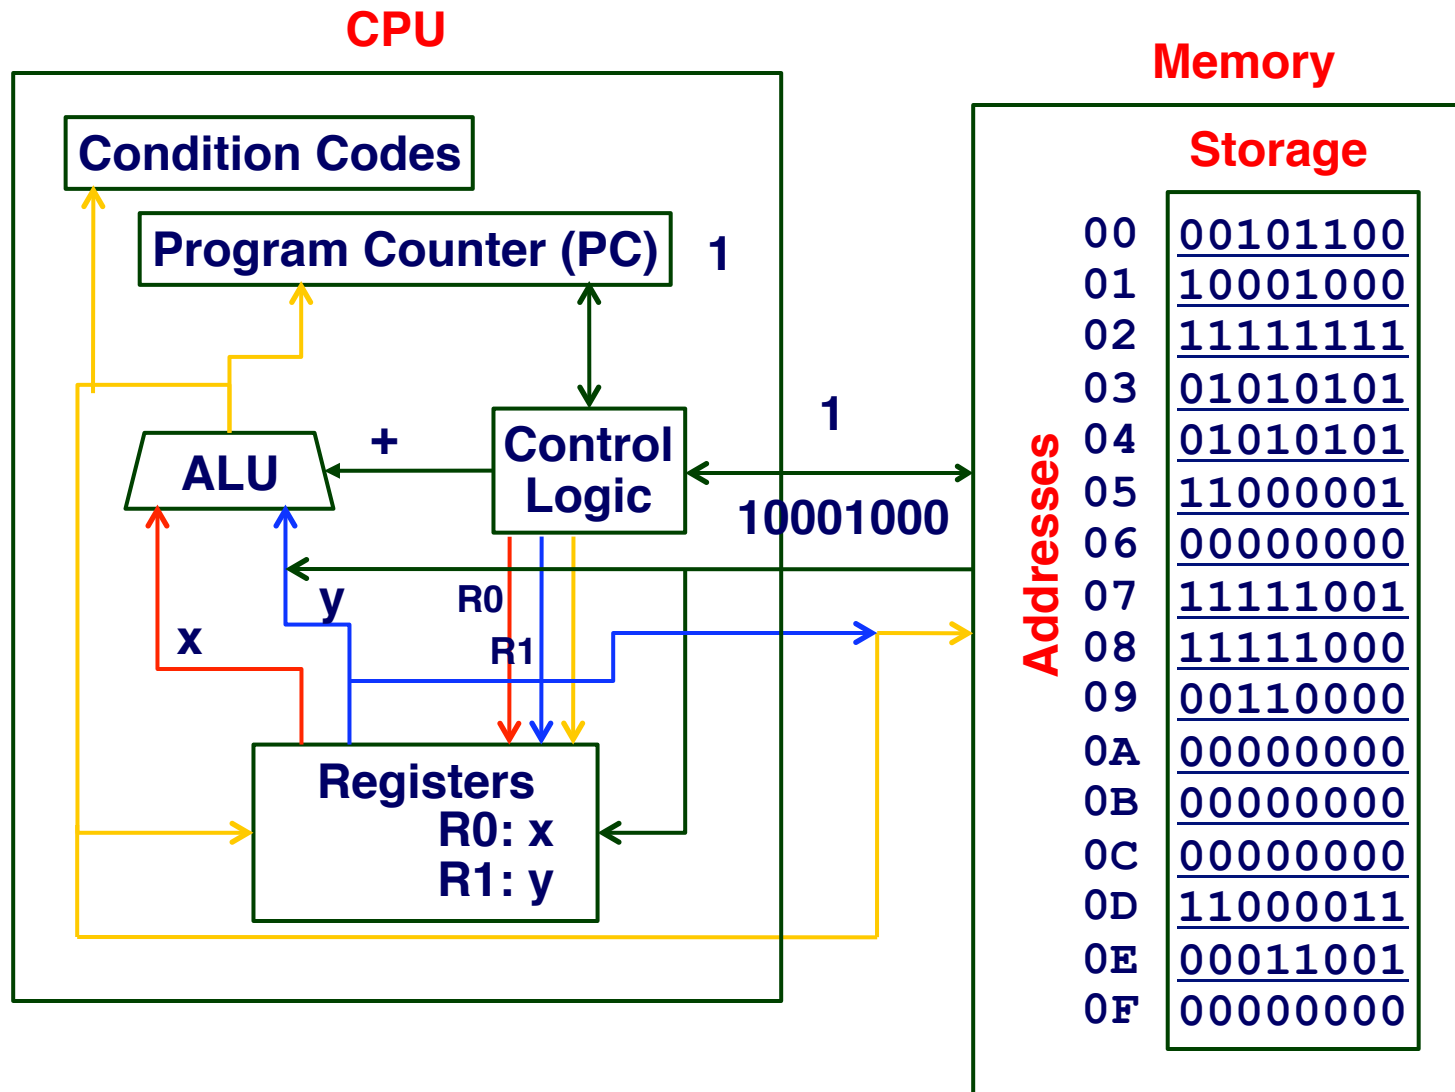
# Putting It All Together



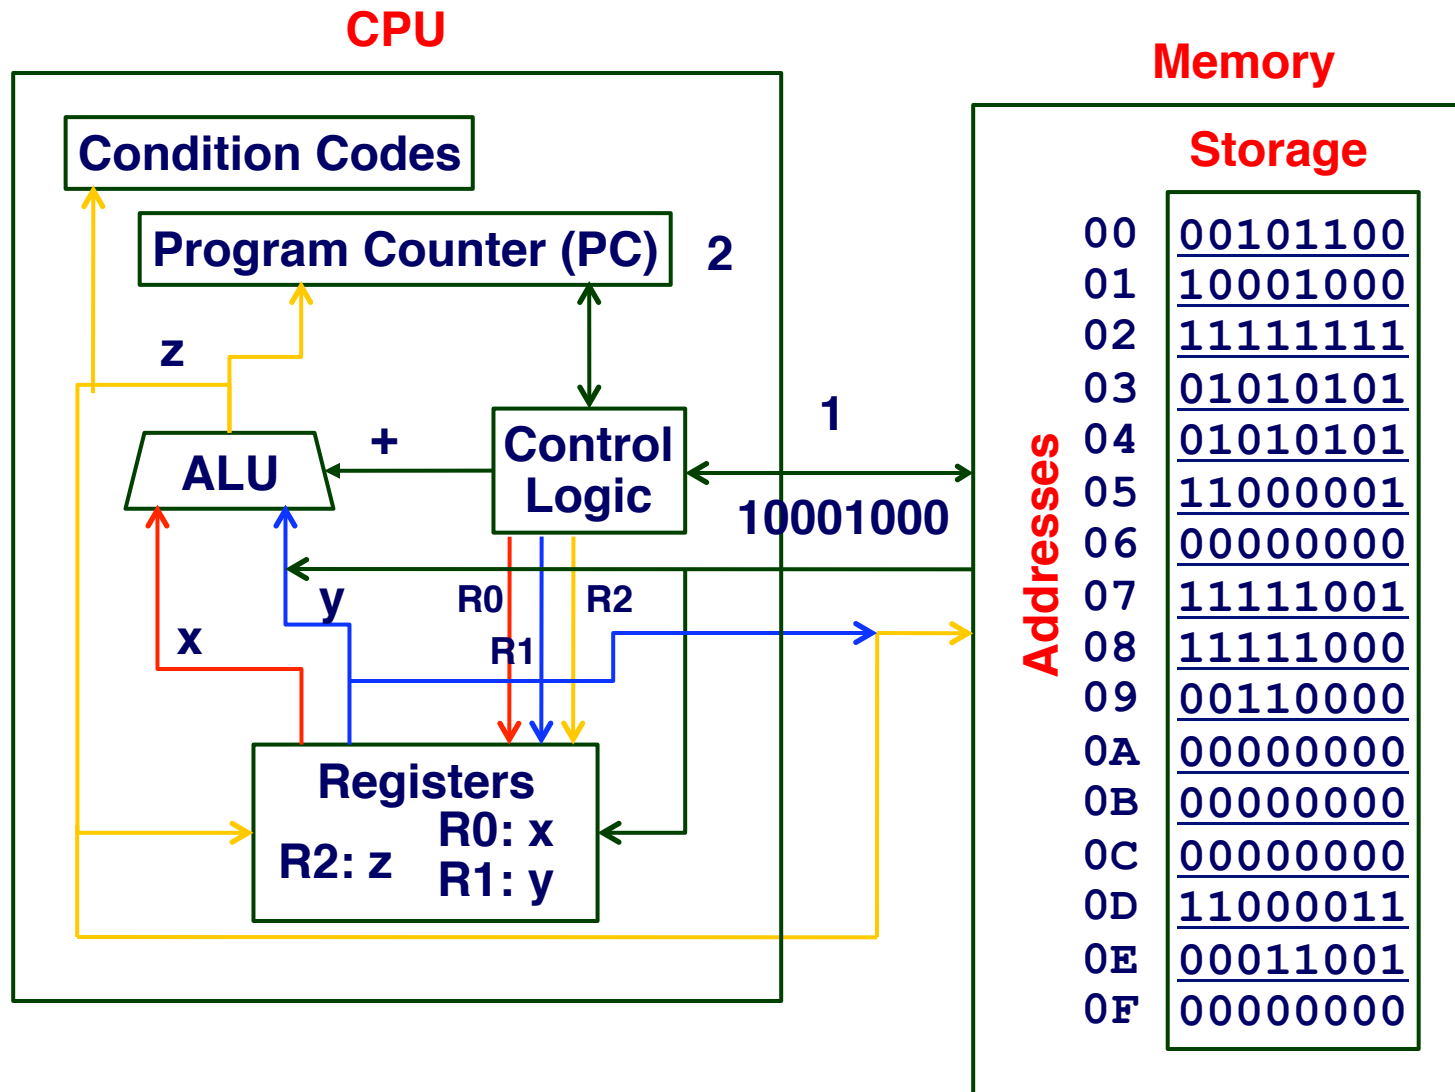
# Putting It All Together



# Putting It All Together



# Putting It All Together



# C & Assembly Code

## Sample C Code

```
int accum;
int sum(int x, int y)
{
    int t = x + y;
    accum += t;
    return t;
}
```

**gcc -O1 -m32 -S code.c**

## Generated Assembly

```
sum:
    push %ebp
    movl %esp, %ebp
    movl 12(%ebp), %eax
    addl 8(%ebp), %eax
    addl %eax, accum
    popl %ebp
    ret
```

# C & Machine Code

## Sample C Code

```
int accum;
int sum(int x, int y){
    int t = x + y;
    accum += t;
    return t;
}
```

**gcc -O1 -m32 -c code.c**

**gdb code.o**

(gdb) x/100xb sum

## objdump -d code.o

```
00000000 <sum>:
0: 55                push  %ebp
1: 89 e5            mov   %esp,%ebp
3: 8b 45 0c         mov   0xc(%ebp),%eax
6: 03 45 08         add   0x8(%ebp),%eax
9: 01 05 00 00 00 00 add   %eax, accum
f: 5d                pop   %ebp
10: c3               ret
```

<sum>: 0x55 0x89 0xe5 0x8b 0x45 0x0c 0x03 0x45

0x8 <sum+8>: 0x08 0x01 0x05 0x00 0x00 0x00 0x00 0x5d

0x10 <sum+16>: 0xc3 Cannot access memory at address 0x11

# Assembly Characteristics

Sequence of simple instructions

Minimal Data Types

- “Integer” data of 1, 2, or 4 bytes
  - Data values
  - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- No aggregate types such as arrays or structures
  - Just contiguously allocated bytes in memory

No type checking

- Interpretation of data format depends on instruction
- No protection against misinterpretation of data



# Assembly Characteristics

## 3 types of Primitive Operations

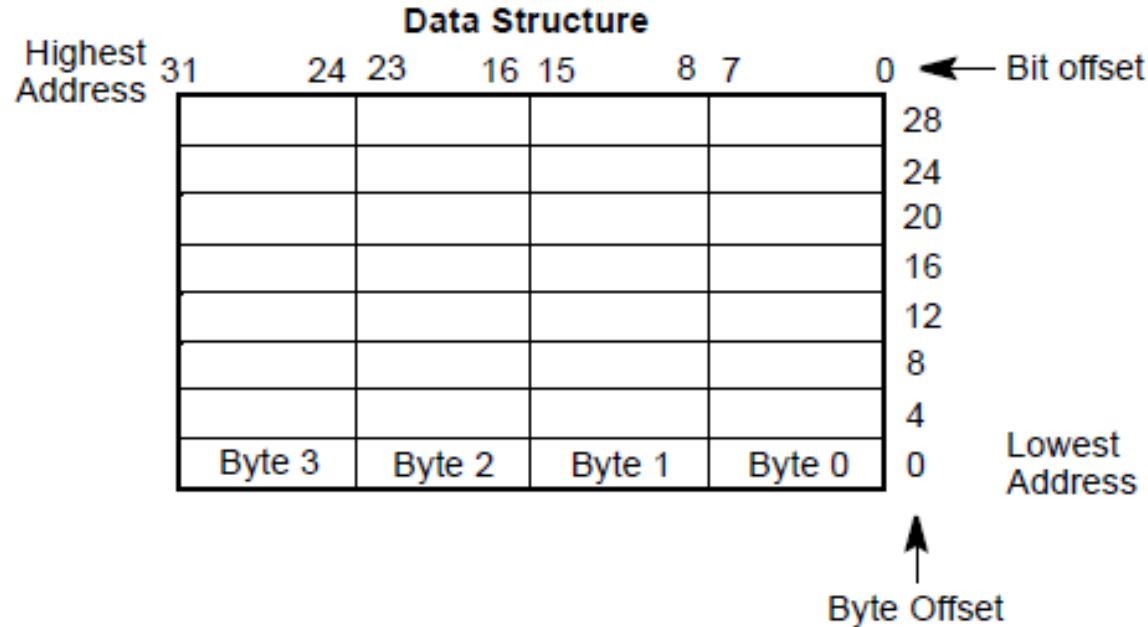
- Perform arithmetic function on register or memory data
- Transfer data between memory and register
  - Load data from memory into register
  - Store register data into memory
- Transfer control
  - Unconditional jumps to/from procedures
  - Conditional branches

# x86 Characteristics

Variable length instructions: 1-15 bytes

Can address memory directly in most instructions

Uses Little-Endian format (Least significant byte in the lowest address)



# Instruction Format

General format:

**opcode operands**

Opcode:

- Short mnemonic for instruction's purpose
  - `movb, addl, etc.`

Operands:

- Immediate, register, or memory
- Number of operands command-dependent

Example:

- `movl %ebx, (%ecx)`

# Machine Representation

Remember, each assembly instruction translated to a sequence of 1-15 bytes

Opcode	addressing mode	other bytes
--------	-----------------	-------------

First, the binary representation of the opcode

Second, instruction specifies the addressing mode

- The type of operands (registers or register and memory)
- How to interpret the operands

Some instructions can be single-byte because operands and addressing mode are implicitly specified by the instruction

- E.g., pushl

# x86 Registers

General purpose registers are 32 bit

- Although operations can access 8-bits or 16-bits portions

Originally categorized into two groups with different functionality

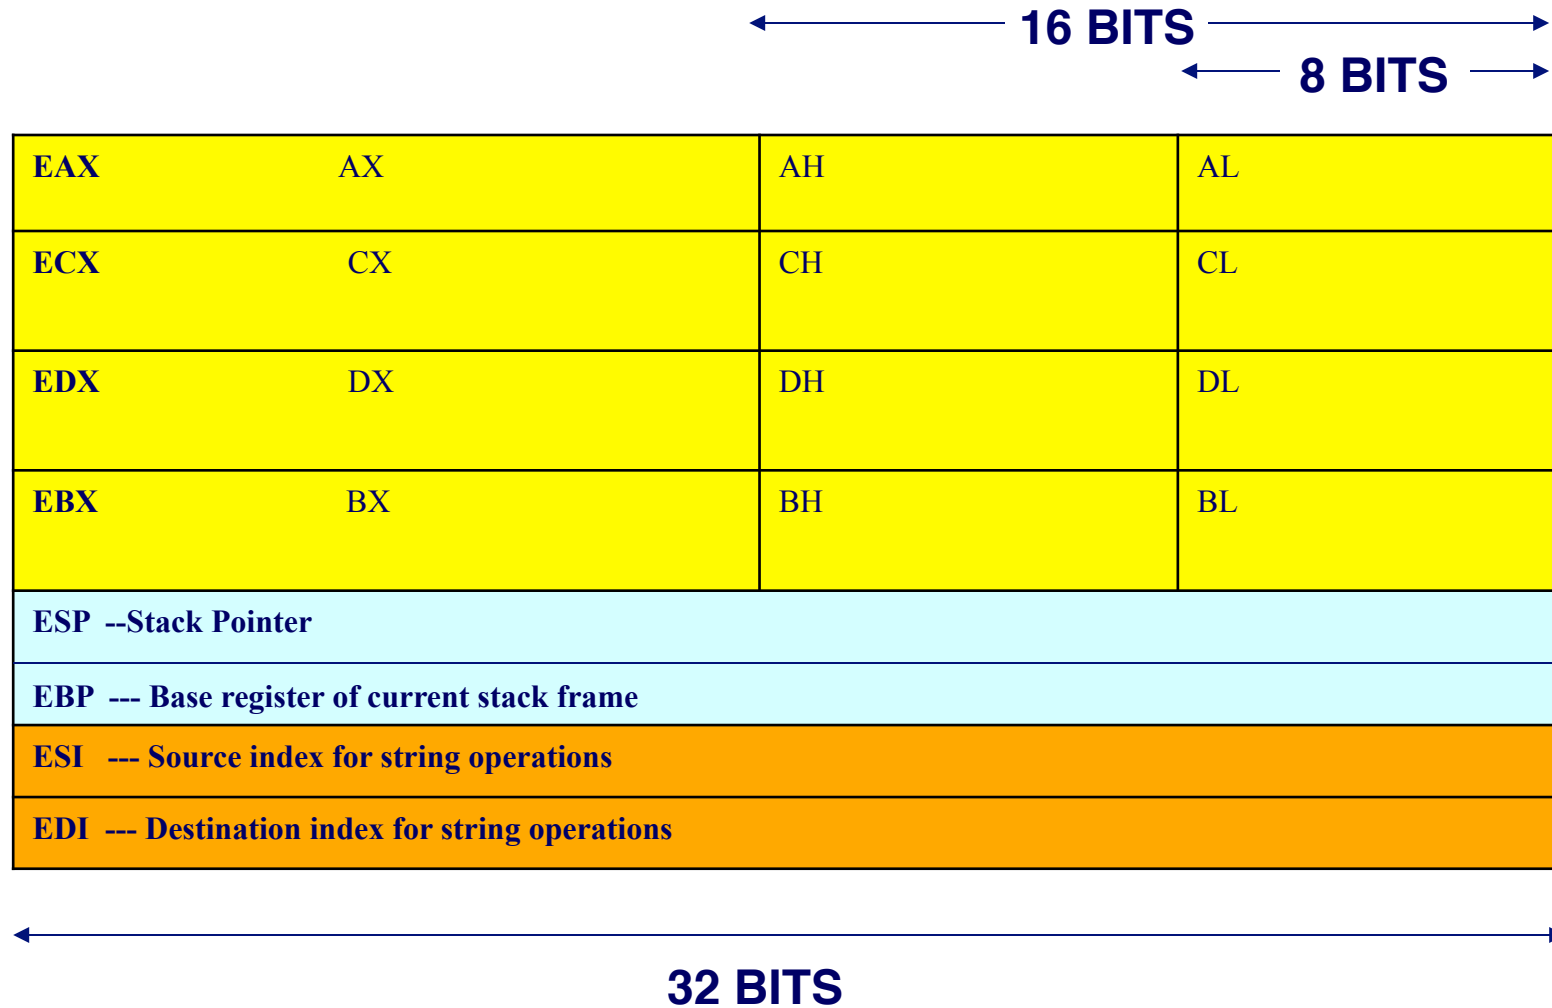
- Data registers (EAX, EBX, ECX, EDX)
  - Holds operands
- Pointer and Index registers (EBP, ESP, EIP,ESI,EDI)
  - Holds references to addresses as well as indexes

Now, the registers are mostly interchangeable

Segment registers

- Holds starting address of program segments
  - CS, DS, SS, ES

# x86 Registers



# x86 Programming

- Mov instructions to move data from/to memory
  - Operands and registers
- Addressing modes
- Understanding swap
- Arithmetic operations
- Condition codes
- Conditional and unconditional branches
- Loops and switch statements

# Data Format

Byte: 8 bits

- E.g., char

Word: 16 bits (2 bytes)

- E.g., short int

Double Word: 32 bits ( 4 bytes)

- E.g., int, float

Quad Word: 64 bits (8 bytes)

- E.g., double

Instructions can operate on any data size

- `movl, movw, movb`
  - **Move double word, word, byte, respectively**
- End character specifies what data size to be used



# MOV instruction

Most common instruction is data transfer instruction

- Mov SRC, DEST: Move source into destination
- SRC and DEST are operands
- DEST is a register or a location
- SRC can be the contents of register, memory location, constant, or a label.
- If you use gcc, you will see `movl <src>, <dest>`
- All the instructions in x86 are 32-bit

Used to copy data:

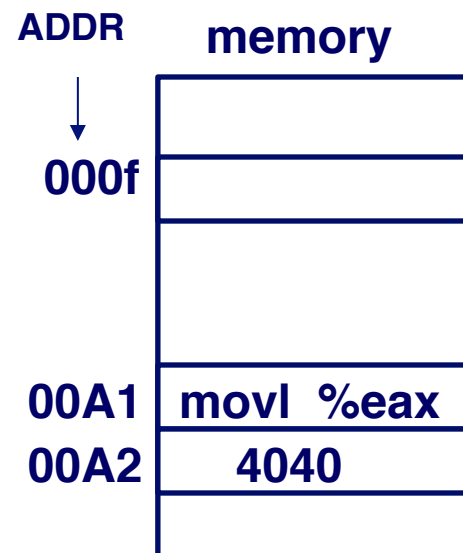
- Constant to register (immediate)
- Memory to register
- Register to memory
- Register to register

**Cannot copy memory to  
memory in a single  
instruction**

# Immediate Addressing

## Operand is immediate

- Operand value is found immediately following the instruction
- Encoded in 1, 2, or 4 bytes
- \$ in front of immediate operand
- E.g., `movl $0x4040, %eax`



# Register Mode Addressing

Use % to denote register

- E.g., %eax

Source operand: use value in specified register

Destination operand: use register as destination for value

Examples:

- `movl %eax, %ebx`
  - Copy content of %eax to %ebx
- `movl $0x4040, %eax` → immediate addressing
  - Copy 0x4040 to %eax
- `movl %eax, 0x0000f` → Absolute addressing
  - Copy content of %eax to memory location 0x0000f

# Indirect Mode Addressing

Content of operand is an address

- Designated as parenthesis around operand

Offset can be specified as immediate mode

Examples:

- `movl (%ebp), %eax`
  - Copy value from memory location whose address is in ebp into eax
- `movl -4(%ebp), %eax`
  - Copy value from memory location whose address is -4 away from content of ebp into eax

# Indexed Mode Addressing

Add content of two registers to get address of operand

- `movl (%ebp, %esi), %eax`
  - Copy value at (address =  $\text{ebp} + \text{esi}$ ) into `eax`
- `movl 8(%ebp, %esi), %eax`
  - Copy value at (address =  $8 + \text{ebp} + \text{esi}$ ) into `eax`

Useful for dealing with arrays

- If you need to walk through the elements of an array
- Use one register to hold base address, one to hold index
  - E.g., implement C array access in a for loop
- Index cannot be `ESP`

# Scaled Indexed Mode Addressing

Multiply the second operand by the scale (1, 2, 4 or 8)

- `movl 0x80(%ebx, %esi, 4), %eax`
  - Copy value at (address =  $\text{ebx} + \text{esi} * 4 + 0x80$ ) into `eax`

Where is it useful?

# Address Computation Examples

<code>%edx</code>	<code>0xf000</code>
<code>%ecx</code>	<code>0x100</code>

Expression	Computation	Address
<code>0x8(%edx)</code>	<code>0xf000 + 0x8</code>	<code>0xf008</code>
<code>(%edx,%ecx)</code>	<code>0xf000 + 0x100</code>	<code>0xf100</code>
<code>(%edx,%ecx,4)</code>	<code>0xf000 + 4*0x100</code>	<code>0xf400</code>
<code>0x80(,%edx,2)</code>	<code>2*0xf000 + 0x80</code>	<code>0x1e080</code>

# movl Operand Combinations

	Source	Destination		C Analog
movl	Imm	Reg	movl \$0x4,%eax	temp = 0x4;
		Mem	movl \$-147, (%eax)	*p = -147;
	Reg	Reg	movl %eax,%edx	temp2 = temp1;
		Mem	movl %eax, (%edx)	*p = temp;
	Mem	Reg	movl (%eax), %edx	temp = *p;

- Cannot do memory-memory transfers with single instruction



# Stack Operations

By convention, `%esp` is used to maintain a stack in memory

- Used to support C function calls

`%esp` contains the address of top of stack

Instructions to push (pop) content onto (off of) the stack

- `pushl %eax`
  - $\text{esp} = \text{esp} - 4$
  - $\text{Memory}[\text{esp}] = \text{eax}$
- `popl %ebx`
  - $\text{ebx} = \text{Memory}[\text{esp}]$
  - $\text{esp} = \text{esp} + 4$

Where does the stack start? We'll discuss later